Connecting to GCP

```
gcloud sql connect db-sp24-demo --user=jingye;

show databases;

use mydb;

show tables;
```

```
^Cjyelin1208@cloudshell:~ (cs411-sp24-demo-419704)$ gcloud sql connect db-sp24-demo --user=jingye
Allowlisting your IP for incoming connection for 5 minutes...done.
Connecting to database with SQL user [jingye].Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 8966
Server version: 8.0.31-google (Google)

Copyright (c) 2000, 2024, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> show databases;
+--------------------+
| Database           |
+--------------------+
| information_schema |
| mydb               |
| mysql              |
| performance_schema |
| sys                |
+--------------------+
5 rows in set (0.01 sec)
```

```
mysql> use mydb;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+----------------+
| Tables_in_mydb |
+----------------+
| CarRank        |
| Cars           |
| FavoriteList   |
| StateTax       |
| Users          |
+----------------+
5 rows in set (0.00 sec)
```

DDL Commands

-CarRank Table

```
CREATE TABLE CarRank(
    CarID DOUBLE,
    Model VARCHAR(50),
    Jan DOUBLE,
    Feb DOUBLE,
    Mar DOUBLE,
    Apr DOUBLE,
    May DOUBLE,
    Jun DOUBLE,
    Jul DOUBLE,
    Aug DOUBLE,
    Sep DOUBLE,
    Oct DOUBLE,
    Nov DOUBLE,
    Dec DOUBLE,
    PRIMARY KEY (CarID ),
    FOREIGN KEY (CarID) REFERENCES Cars(CarID)
);
```

-Cars Table

```
CREATE TABLE Cars(
CarID DOUBLE,
CarName VARCHAR(10),
Year DOUBLE,
SellingPrice DOUBLE,
Fuel VARCHAR(10),
Transmission VARCHAR(10),
PRIMARY KEY (userID),
FOREIGN KEY (stateID) REFERENCES StateTax(stateID)
);
```

-FavoriteList Table

```
CREATE TABLE FavoriteList(
ListID DOUBLE,
CarID DOUBLE,
UserID DOUBLE,
PRIMARY KEY (ListID),
FOREIGN KEY (CarID) REFERENCES Cars(CarID),
FOREIGN KEY (UserID) REFERENCES Users(UserID)
);
```

-StateTax Table

```
CREATE TABLE StateTax(
State VARCHAR(50),
StateID DOUBLE,
TaxRate DOUBLE,
PRIMARY KEY (StateID),
FOREIGN KEY (StateID) REFERENCES Users(StateID)
);
```

-Users Table

```
CREATE TABLE Users(
userID DOUBLE,
userName VARCHAR(10),
gender VARCHAR(10),
stateID DOUBLE,
PRIMARY KEY (userID),
FOREIGN KEY (StateID) REFERENCES StateTax(StateID)
);
```

Inserting Data

```
mysql> SELECT COUNT(*) FROM Users;
+----------+
| COUNT(*) |
+----------+
|     1203 |
+----------+
1 row in set (0.03 sec)

mysql> SELECT COUNT(*) FROM FavoriteList;
+----------+
| COUNT(*) |
+----------+
|     1100 |
+----------+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM Cars;
+----------+
| COUNT(*) |
+----------+
|     4340 |
+----------+
1 row in set (0.00 sec)
```

Advanced Queries

1.
Advanced features: Joining Multiple Relations & Conditional Filtering and Ordering
This SQL query is designed to retrieve a list of the top 15 favorite cars more expensive than 2000 and later than 2010 based on their sales in December.

```
SELECT
   u.UserName,
   c.CarName,
   cr.Dece AS `Sold`,
   'December' AS `Month`
FROM
   Users u
   JOIN FavoriteList f ON u.UserId = f.UserID
   JOIN Cars c ON f.CarID = c.CarID
   JOIN CarRank cr ON c.CarID = cr.CarID
WHERE
   cr.Dece > 0 AND c.SellingPrice>2000 AND Year>2010
ORDER BY
   cr.Dece DESC
LIMIT 15;
```

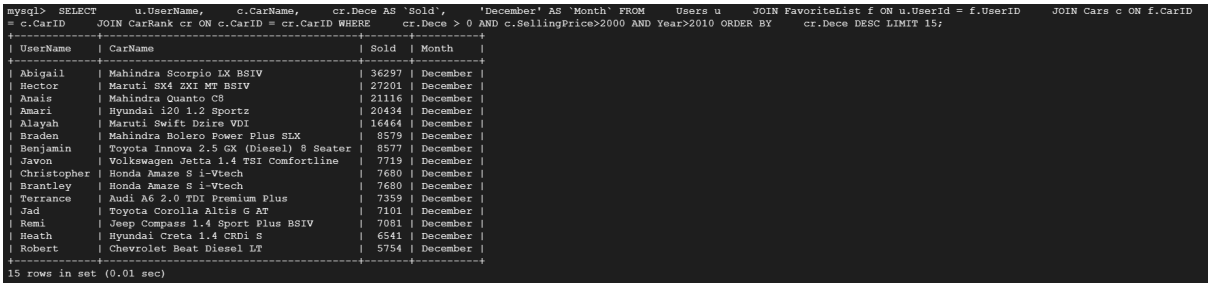

```
mysql> SELECT      u.UserName,     c.CarName,     cr.Dece AS `Sold`,     'December' AS `Month` FROM     Users u     JOIN FavoriteList f ON u.UserId = f.UserID     JOIN Cars c ON f.CarID
= c.CarID     JOIN CarRank cr ON c.CarID = cr.CarID WHERE     cr.Dece > 0 AND c.SellingPrice>2000 AND Year>2010 ORDER BY     cr.Dece DESC LIMIT 15;
+-------------+-----------------------------------------+-------+----------+
| UserName    | CarName                                 | Sold  | Month    |
+-------------+-----------------------------------------+-------+----------+
| Abigail     | Mahindra Scorpio LX BSIV                | 36297 | December |
| Hector      | Maruti SX4 ZXI MT BSIV                  | 27201 | December |
| Anais       | Mahindra Quanto C8                      | 21116 | December |
| Amari       | Hyundai i20 1.2 Sportz                  | 20434 | December |
| Alayah      | Maruti Swift Dzire VDI                  | 16464 | December |
| Braden      | Mahindra Bolero Power Plus SLX          |  8579 | December |
| Benjamin    | Toyota Innova 2.5 GX (Diesel) 8 Seater |  8577 | December |
| Javon       | Volkswagen Jetta 1.4 TSI Comfortline    |  7719 | December |
| Christopher | Honda Amaze S i-Vtech                   |  7680 | December |
| Brantley    | Honda Amaze S i-Vtech                   |  7680 | December |
| Terrance    | Audi A6 2.0 TDI Premium Plus            |  7359 | December |
| Jad         | Toyota Corolla Altis G AT               |  7101 | December |
| Remi        | Jeep Compass 1.4 Sport Plus BSIV        |  7081 | December |
| Heath       | Hyundai Creta 1.4 CRDi S                |  6541 | December |
| Robert      | Chevrolet Beat Diesel LT                |  5754 | December |
+-------------+-----------------------------------------+-------+----------+
15 rows in set (0.01 sec)
```

2.
Advanced features: Joining Multiple Relations & Aggregation with GROUP BY & Subqueries
This query is to calculate the average selling price of vehicles that are more expensive than 100,000 (manual transmission or automatic transmission) preferred by users of different genders.

```
SELECT
  u.gender,
  CASE
    WHEN c.Transmission = 'Manual' THEN 'Manual'
    WHEN c.Transmission = 'Automatic' THEN 'Automatic'
    ELSE 'Other'
  END AS transmission_type,
  AVG(c.SellingPrice) AS avg_selling_price
```

```
FROM
  FavoriteList f
JOIN
  Cars c ON f.CarID = c.CarID
JOIN
  Users u ON f.userID = u.userID
GROUP BY
  u.gender,
  Transmission_type
Having avg_selling_price > 100000;
```

```
+--------+-------------------+--------------------+
| gender | transmission_type | avg_selling_price |
+--------+-------------------+--------------------+
| Female | Manual            | 396777.9862745098 |
| Male   | Manual            | 401168.1821862348 |
| Male   | Automatic         |        1425539.98 |
| Female | Automatic         |        1373459.94 |
+--------+-------------------+--------------------+
4 rows in set (0.01 sec)
```

3. Advanced features: Aggregation via Group By & Joining Multiple Relations
This query combines information from three tables to find the most popular cars by state.

```
SELECT
   StateTax.State,
   Cars.CarName,
   COUNT(FavoriteList.CarID) AS PopularityScore
FROM
   FavoriteList
JOIN
   Users ON FavoriteList.UserID = Users.userID
JOIN
   StateTax ON Users.stateID = StateTax.StateID
JOIN
   Cars ON FavoriteList.CarID = Cars.CarID
WHERE Cars.Year  > 2010
GROUP BY
   StateTax.State, Cars.CarName
ORDER BY
   StateTax.State,  PopularityScore DESC
```

Limit 15;

State	CarName	PopularityScore
Alabama	Tata Indica Vista Quadrajet LS	1
Alabama	Hyundai i10 Sportz AT	1
Alabama	Hyundai Creta 1.6 CRDi SX	1
Alabama	Fiat Linea 1.3 Emotion	1
Alabama	Mahindra XUV500 AT W10 FWD	1
Alabama	Hyundai Grand i10 Sportz	1
Alabama	Hyundai Verna 1.6 CRDi AT SX	1
Alabama	Tata Tiago 1.05 Revotorq XE	1
Alabama	Hyundai i20 Asta 1.4 CRDi	1
Alabama	Maruti Wagon R VXI BS IV	1
Alabama	Mahindra Scorpio S7 140 BSIV	1
Alabama	Maruti Alto 800 LXI	1
Alabama	Mahindra Quanto C4	1
Alaska	Maruti Vitara Brezza ZDi Plus	1
Alaska	Renault KWID RXT Optional	1

15 rows in set (0.00 sec)

4. Advanced features: Aggregation via Group By & AVG Function & Having Condition & Order By
This query combines information from three tables to find the most popular cars by state.

```
SELECT
   Year,
   Fuel,
   AVG(SellingPrice) AS AverageSellingPrice
FROM
   Cars
GROUP BY
   Year, Fuel
HAVING
   AVG(SellingPrice) > (SELECT AVG(SellingPrice) FROM Cars WHERE Year <
YEAR(CURRENT_DATE()))
ORDER BY
   Year DESC, AverageSellingPrice DESC;
```

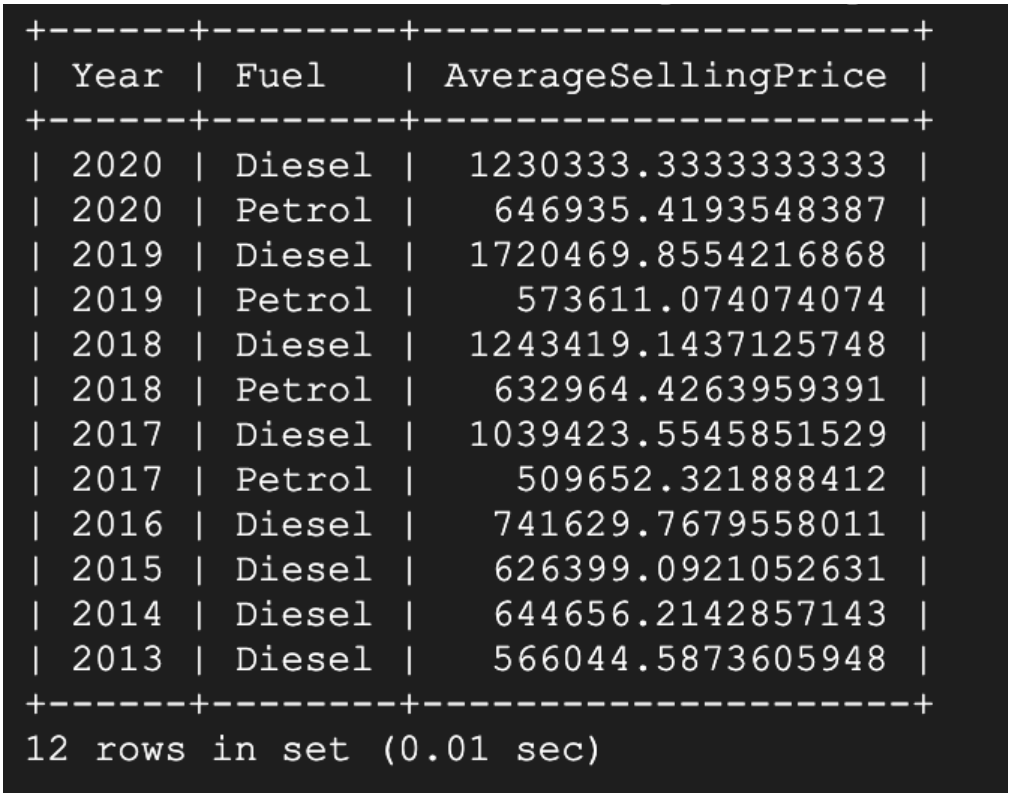

Year	Fuel	AverageSellingPrice
2020	Diesel	1230333.3333333333
2020	Petrol	646935.4193548387
2019	Diesel	1720469.8554216868
2019	Petrol	573611.074074074
2018	Diesel	1243419.1437125748
2018	Petrol	632964.4263959391
2017	Diesel	1039423.5545851529
2017	Petrol	509652.321888412
2016	Diesel	741629.7679558011
2015	Diesel	626399.0921052631
2014	Diesel	644656.2142857143
2013	Diesel	566044.5873605948

12 rows in set (0.01 sec)

5. Advanced features: SUM function and Ordering
This query displays car details along with their total sales ranking score from January to December.

```
SELECT
   Cars.CarID,
   Cars.CarName,
   Cars.Year,
   Cars.SellingPrice,
   Cars.Fuel,
   Cars.Transmission,
   (SELECT SUM(Jan + Feb + Mar + Apr + May + Jun + Jul + Aug + Sep + Oct + Nov + Dece)
    FROM CarRank
    WHERE CarRank.CarID = Cars.CarID) AS TotalRankScore
FROM
   Cars
ORDER BY TotalRankScore DESC
LIMIT 15;
```

6. Advanced features: Joining Multiple Relations & Conditional filtering

This query shows detailed car information in a user's favorite list, including how much tax they'd pay based on their state.

```
SELECT
   Users.userName,
   Cars.CarName,
   Cars.Year,
   Cars.SellingPrice,
   StateTax.TaxRate,
   (Cars.SellingPrice * (StateTax.TaxRate / 100)) AS TaxAmount
FROM
   FavoriteList
JOIN
   Cars ON FavoriteList.CarID = Cars.CarID
JOIN
   Users ON FavoriteList.UserID = Users.userID
JOIN
   StateTax ON Users.stateID = StateTax.StateID
WHERE
   FavoriteList.UserID = 311
Limit 15;
```

Indexing (for the first four queries)

Query 1

```
mysql> EXPLAIN ANALYZE
    -> SELECT
    ->     u.UserName,
    ->     c.CarName,
    ->     cr.Dece AS `Sold`,
    ->     'December' AS `Month`
    -> FROM
    ->     Users u
    ->     JOIN FavoriteList f ON u.UserId = f.UserID
    ->     JOIN Cars c ON f.CarID = c.CarID
    ->     JOIN CarRank cr ON c.CarID = cr.CarID
    -> WHERE
    ->     cr.Dece > 0 AND c.SellingPrice>2000 AND Year>2010
    -> ORDER BY
    ->     cr.Dece DESC
    -> LIMIT 15;
```

This query has a high cost due to our use of a set operation

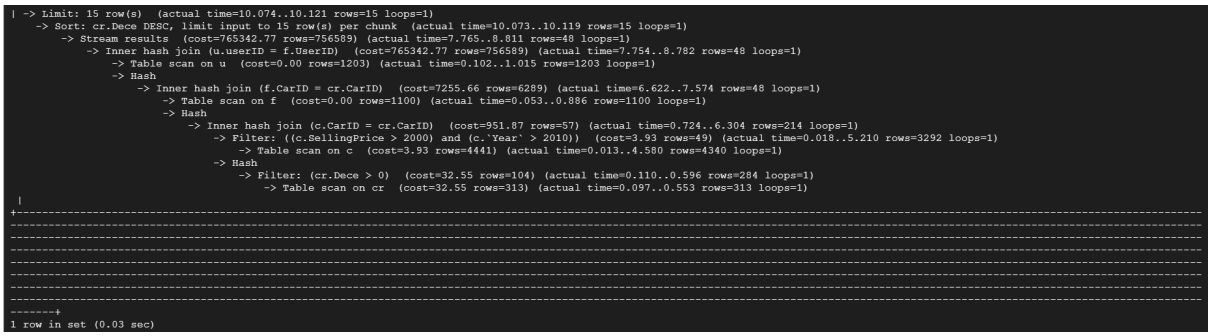

Attempt1: CREATE INDEX add_c_Dece_idx ON CarRank(Dece)

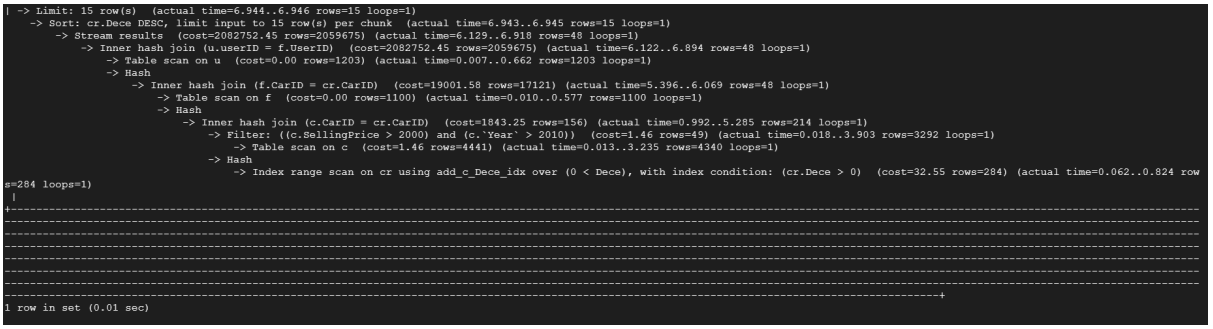

This design of the index is able to reduce the time. After utilizing an index scan on the joined table, the query efficiently located relevant data rows. This reduction in the number of times the table needed to be scanned led to a significant improvement in query performance. It reduces the time of scanning the table CarRank from 0.097 to 0.062. And also reduces the general cost in the first line.

Attempt2: CREATE INDEX add_c_sellprice_idx ON Cars(SellingPrice)

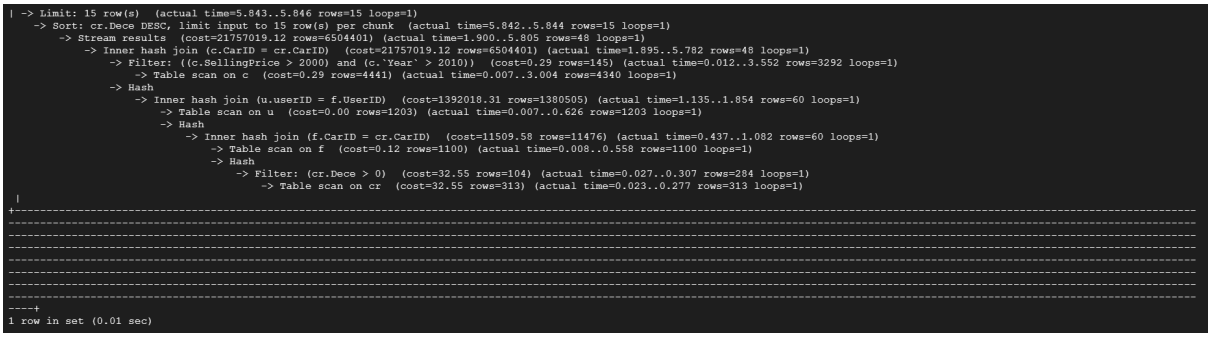

This index design effectively reduces query execution time. By utilizing an index scan on the joined table, the query efficiently locates relevant data rows. This reduction in the number of times the table needs to be scanned significantly improves query performance. It also minimizes the overall cost, as mentioned in the first line.

Attempt3: CREATE INDEX add_c_year_idx ON Cars(Year)

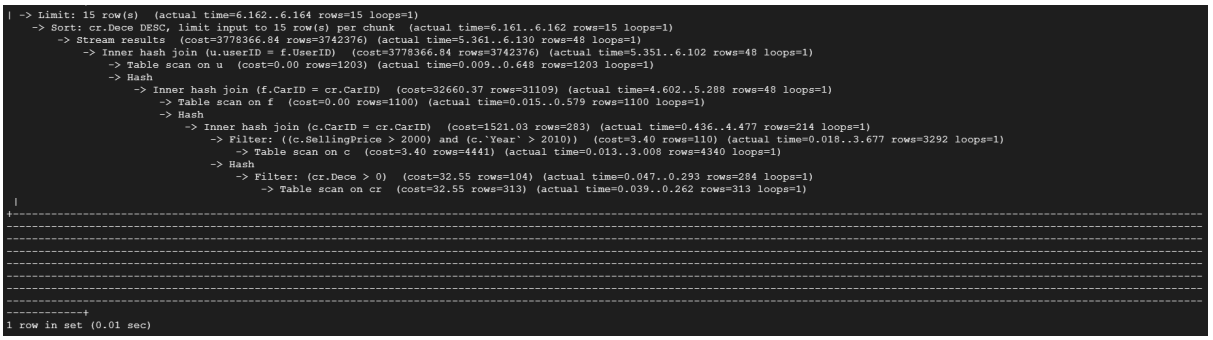

This index design significantly reduces query execution time. By employing an index scan on the joined table, the query efficiently identifies relevant data rows. Consequently, the reduced number of table scans results in a substantial improvement in query performance, leading to cost savings overall.

Query 2

```
mysql> EXPLAIN ANALYZE
    -> SELECT
    ->   u.gender,
    ->   CASE
    ->     WHEN c.Transmission = 'Manual' THEN 'Manual'
    ->     WHEN c.Transmission = 'Automatic' THEN 'Automatic'
    ->     ELSE 'Other'
    ->   END AS transmission_type,
    ->   AVG(c.SellingPrice) AS avg_selling_price
    -> FROM
    ->   FavoriteList f
    -> JOIN
    ->   Cars c ON f.CarID = c.CarID
    -> JOIN
    ->   Users u ON f.userID = u.userID
    -> GROUP BY
    ->   u.gender,
    ->   Transmission_type
    -> Having avg_selling_price > 100000;
```

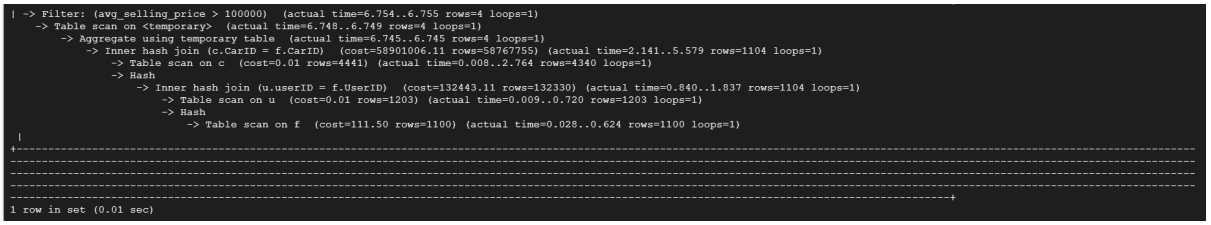

Attempt1: CREATE INDEX add_u_genderid_idx ON Users(gender)

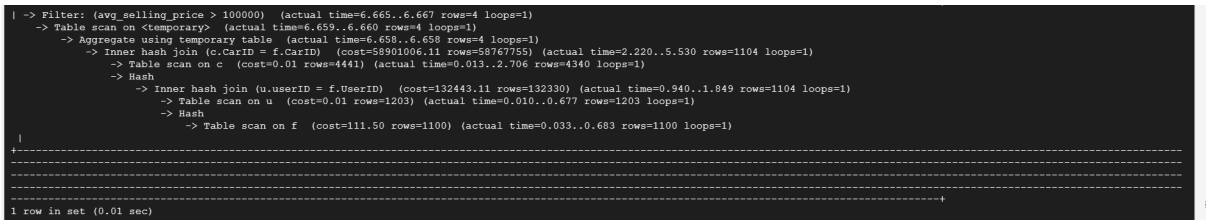

There is no significant change in actual time of all and that of aggregate. This might be because gender itself has few different values, making the index no sense. And group by need to scan the table again.

Attempt2: CREATE INDEX add_c_transmission_idx ON Cars(Transmission)

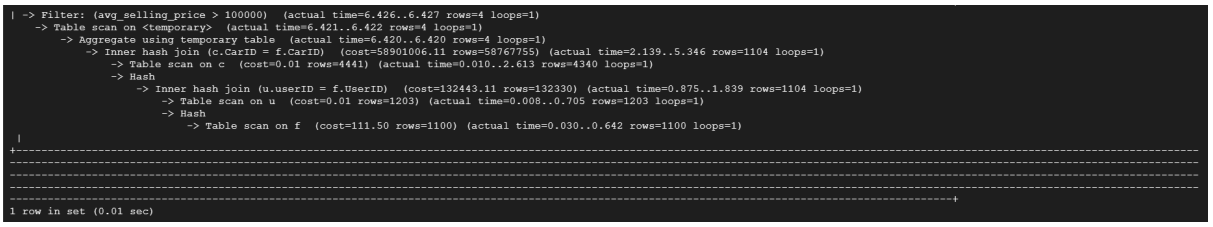

There is no significant change in actual time of all and that of aggregate. This might be because transmission itself has few different values, making the index no sense. And group by need to scan the table again.

Attempt3: CREATE INDEX add_c_sellprice_idx ON Cars(SellingPrice)

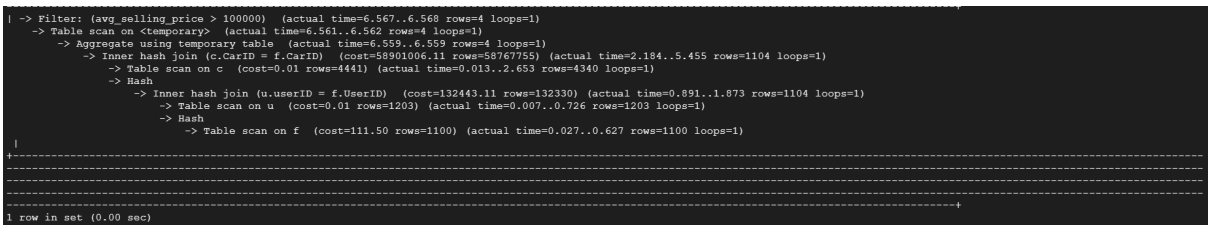

There is no significant change in actual time of all and that of aggregate. This might be because the average selling price will be calculated again, which means the origin value should be scanned from the table.

Query 3

```
mysql> EXPLAIN ANALYZE
    -> SELECT
    ->     StateTax.State,
    ->     Cars.CarName,
    ->     COUNT(FavoriteList.CarID) AS PopularityScore
    -> FROM
    ->     FavoriteList
    -> JOIN
    ->     Users ON FavoriteList.UserID = Users.userID
    -> JOIN
    ->     StateTax ON Users.stateID = StateTax.StateID
    -> JOIN
    ->     Cars ON FavoriteList.CarID = Cars.CarID
    -> WHERE Cars.Year  > 2010
    -> GROUP BY
    ->     StateTax.State, Cars.CarName
    -> ORDER BY
    ->     StateTax.State,  PopularityScore DESC
    -> Limit 15;
```

```
| -> Limit: 15 row(s)  (actual time=8.041..8.043 rows=15 loops=1)
    -> Sort: StateTax.State, PopularityScore DESC, limit input to 15 row(s) per chunk  (actual time=8.040..8.041 rows=15 loops=1)
        -> Table scan on <temporary>  (actual time=7.668..7.803 rows=783 loops=1)
            -> Aggregate using temporary table  (actual time=7.666..7.666 rows=783 loops=1)
                -> Inner hash join (Cars.CarID = FavoriteList.CarID)  (cost=10861244.53 rows=3329507) (actual time=2.947..6.588 rows=804 loops=1)
                    -> Filter: (Cars.`Year` > 2010)  (cost=0.28 rows=148) (actual time=0.014..2.982 rows=3292 loops=1)
                        -> Table scan on Cars  (cost=0.28 rows=4441) (actual time=0.010..2.607 rows=4340 loops=1)
                    -> Hash
                        -> Inner hash join (FavoriteList.UserID = Users.userID)  (cost=681034.69 rows=674883) (actual time=1.655..2.582 rows=1063 loops=1)
                            -> Table scan on FavoriteList  (cost=0.00 rows=1100) (actual time=0.011..0.627 rows=1100 loops=1)
                            -> Hash
                                -> Inner hash join (Users.stateID = StateTax.StateID)  (cost=6142.21 rows=6135) (actual time=0.150..1.282 rows=1155 loops=1)
                                    -> Table scan on Users  (cost=0.27 rows=1203) (actual time=0.012..0.812 rows=1203 loops=1)
                                    -> Hash
                                        -> Table scan on StateTax  (cost=5.35 rows=51) (actual time=0.029..0.104 rows=51 loops=1)
 |
1 row in set (0.01 sec)
```

Attempt1: CREATE INDEX add_c_year_idx ON Cars(Year)

```
| -> Limit: 15 row(s)  (actual time=7.891..7.893 rows=15 loops=1)
    -> Sort: StateTax.State, PopularityScore DESC, limit input to 15 row(s) per chunk  (actual time=7.890..7.891 rows=15 loops=1)
        -> Table scan on <temporary>  (actual time=7.584..7.695 rows=783 loops=1)
            -> Aggregate using temporary table  (actual time=7.582..7.582 rows=783 loops=1)
                -> Inner hash join (Cars.CarID = FavoriteList.CarID)  (cost=22974880.30 rows=16469007) (actual time=2.687..6.397 rows=804 loops=1)
                    -> Filter: (Cars.`Year` > 2010)  (cost=0.11 rows=329) (actual time=0.015..3.076 rows=3292 loops=1)
                        -> Table scan on Cars  (cost=0.11 rows=4441) (actual time=0.011..2.696 rows=4340 loops=1)
                    -> Hash
                        -> Inner hash join (FavoriteList.UserID = Users.userID)  (cost=681034.69 rows=674883) (actual time=1.351..2.285 rows=1063 loops=1)
                            -> Table scan on FavoriteList  (cost=0.00 rows=1100) (actual time=0.007..0.651 rows=1100 loops=1)
                            -> Hash
                                -> Inner hash join (Users.stateID = StateTax.StateID)  (cost=6142.21 rows=6135) (actual time=0.083..1.017 rows=1155 loops=1)
                                    -> Table scan on Users  (cost=0.27 rows=1203) (actual time=0.009..0.691 rows=1203 loops=1)
                                    -> Hash
                                        -> Table scan on StateTax  (cost=5.35 rows=51) (actual time=0.021..0.050 rows=51 loops=1)
 |
1 row in set (0.01 sec)
```

This design of the index is able to reduce the time. After utilizing an index scan on the joined table, the query efficiently located relevant data rows. This reduction in the number of times the table needed to be scanned led to a significant improvement in query performance. It reduces the general cost in the first line.

Attempt2: CREATE INDEX add_st_State_idx ON StateTax(State)

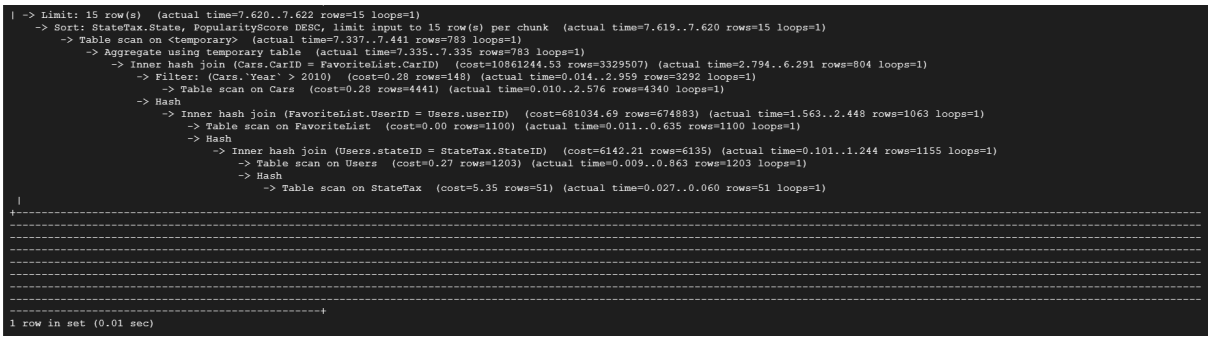

This index design significantly reduces query execution time. By employing an index scan on the joined table, the query efficiently identifies relevant data rows. Consequently, the reduced number of table scans results in a substantial improvement in query performance, leading to cost savings overall.

Attempt3: CREATE INDEX add_c_CarName_idx ON Cars(CarName)

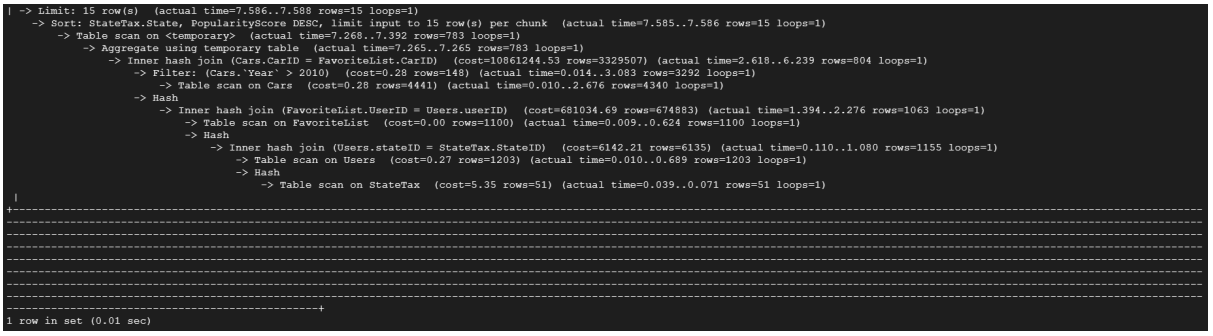

This index design of car name in cars table is similar with car year, which reduced number of table scans.

Query 4

```
mysql> EXPLAIN ANALYZE
    -> SELECT
    ->     Year,
    ->     Fuel,
    ->     AVG(SellingPrice) AS AverageSellingPrice
    -> FROM
    ->     Cars
    -> GROUP BY
    ->     Year, Fuel
    -> HAVING
    ->     AVG(SellingPrice) > (SELECT AVG(SellingPrice) FROM Cars WHERE Year < YEAR(CURRENT_DATE()))
    -> ORDER BY
    ->     Year DESC, AverageSellingPrice DESC;
```

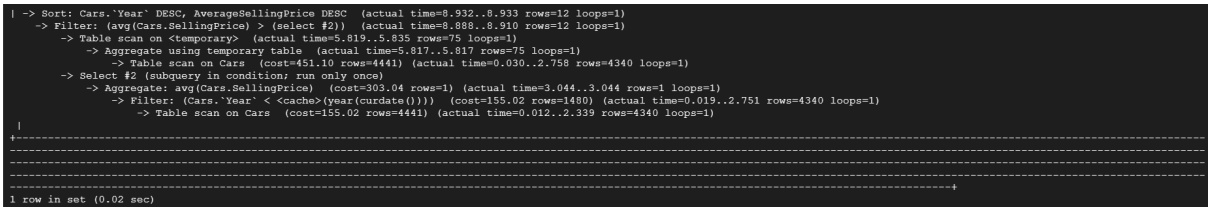

Attempt1: CREATE INDEX add_c_fuel_idx ON Cars(Fuel)

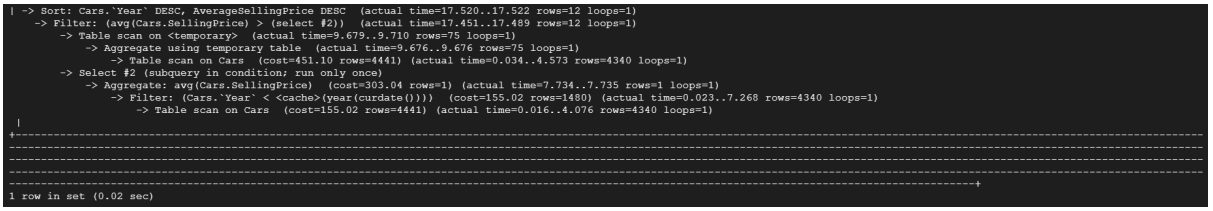

This design increased the total actual time. Since the cars table still needs to be scanned with year and price, this design makes it more complicated and costs more time.

Attempt2: CREATE INDEX add_c_year_idx ON Cars(Year)

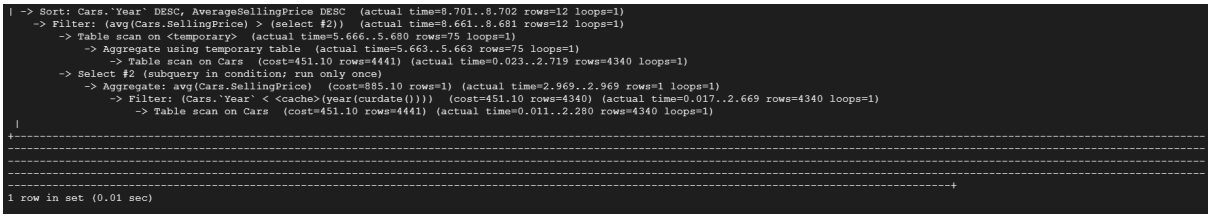

This design does not make improvements. This might be because the table scan cannot be reduced when using the car year index.

Attempt3: CREATE INDEX add_c_sellprice_idx ON Cars(SellingPrice)

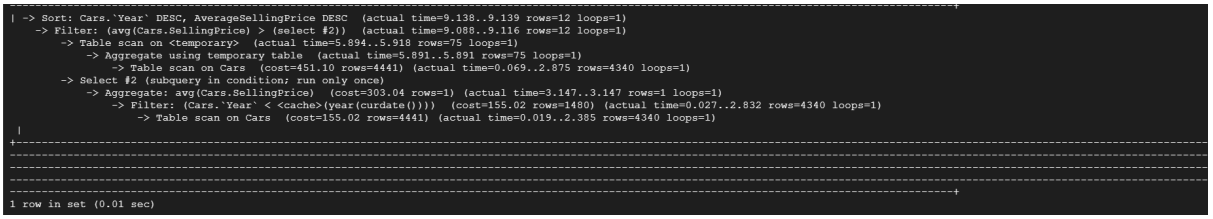

This design of using sell price as index does not make any improvement on the actual time. It does not reduce the time of aggregate process and also no reduce of table scan.